

Converged IT Infrastructure in Life Sciences

Org Challenges, Networks, and Science DMZs, Oh My!

We've noticed a few things

We have a unique perspective across much of life sciences

- ▶ **Big Data has arrived in Life Sciences**
- ▶ **Data is being generated at unprecedented rates**
- ▶ **Research and Biomedical Orgs were caught off guard**
- ▶ **IT running to catch up, limited budgets**
- ▶ **Money is tight, Orgs reluctant to invest in Bio-IT**

25% of all Life Scientists will require HPC in 2015!

Research is hard


It's being made harder by lack of support and infrastructure that works for them

- ▶ **Scientists are getting frustrated**
- ▶ **Stubborn, fight for what they need**
- ▶ **They will build a cluster under their desk if it gets the job done**
- ▶ **In general they will win against IT, eventually**



Big Picture / Meta Issue

- ▶ **HUGE revolution in the rate at which lab platforms are being redesigned, improved & refreshed**
- ▶ **IT not a part of the conversation, running to catch up**




The Central Problem Is ...

Science progressing way faster than IT can refresh/
change

- ▶ **Instrumentation & protocols are changing FAR FASTER than we can refresh our Research-IT & Scientific Computing infrastructure**
 - Bench science is changing month-to-month ...
 - ... while our IT infrastructure only gets refreshed every 2-7 years
- ▶ **We have to design systems TODAY that can support unknown research requirements & workflows over many years (gulp ...)**

The Central Problem Is ...

- ▶ **The easy period is over**
- ▶ **5 years ago - under the desk solutions worked**
- ▶ **This doesn't work any more!**





The new normal for informatics





What are the drivers in Bio-IT today?




Genomics: Next Generation Sequencing (NGS)



The big deal about DNA



It's like the hard drive of life

- ▶ **DNA is the template of life**
- ▶ **DNA is read --> RNA**
- ▶ **RNA is read --> Proteins**
- ▶ **Proteins are the functional machinery that make life possible**
- ▶ **Understanding the template = understanding basis for disease**



How does NGS work?

Sequencing by Synthesis




A sequence chromatogram showing DNA sequence data. The sequence reads:

```

ACACTGGAACTGAGACACGGTCCAGACTC
GGGGAAACCCCTGAAGCAGCAACGCCGGT
TTAGGGAAAGAACCATGACGGTACCTACG
GTAATACGGAGGGTGCAGCGTTACTCGG
AAGTCTTTGTGA MUTATION AATCTA
ATCTAGAGTGAGGGAGAGGGCAGATGGAA
GGAAATACCCATTGCGAAGGGCAGATCTGCT
GAGCAAACAGGATTAGATAACCCCTGGTAGT
CTAGTCACGGCAGTAATGCACCTAACGG

```

The word "MUTATION" is highlighted in blue, indicating a sequencing error or artifact.

How does NGS work?


Reference assembly, variant calling



The Human Genome

Gateway to personalized medicine

- ▶ **3.2 Gbp**
- ▶ **23 chromosomes**
- ▶ **~21,000 genes**
- ▶ **Over 55M known variations**



The Problem...

...and why NGS is the primary driver

- ▶ **Sequencers are now relatively cheap and fast**
- ▶ **Some can generate a human genome in 18 hours, for \$2,000**
- ▶ **Everyone is doing it**
- ▶ **Can generate 3TB of data in that time**
- ▶ **First genome took 13 years and \$2.7B to complete**



A yellow school bus is parked on a grassy field. The bus has "PROVIDENCE HILL SCHOOL DISTRICT" printed on its side. In the background, there are trees and a chain-link fence.

Other Methodologies Not Far Behind

High-throughput Imaging


- ▶ **Robotics screening millions of compounds on live cells 24/7**

- Not as much data as genomics in volume, but just as complex
- Data volumes in the 10's TB/week



- ▶ **Confocal Imaging**


- Scanning 100's of tissue sections/week, each with 10's of scans, each with 20-40 layers and multiple fluorescent channels
- Data volumes in the 1's - 10's TB/week



High-res medical imaging

High-power, dense detector MRI scanners in use 24/7 at large research hospitals

- ▶ **Creating 3D models of brains, comparing large datasets**
- ▶ **Using those models to perform detailed neurosurgery with real-time analytic feedback from supercomputer in the OR (cool stuff)**
- ▶ **Also generates 10's of TB/week**



This is a huge problem

- ▶ Causing a literal deluge of data, in the 10's of Petabytes
- ▶ NIH generating 1.5PB of data/month
- ▶ First real case in life science where 100Gb networking might really be needed
- ▶ But, not enough storage or compute





Most common high-speed network: FedEx

A black hole at the center of a galaxy, shown from a low-angle perspective. A bright, multi-colored accretion disk surrounds the black hole, with light bending around it. Two powerful, blue-colored jets of matter are ejected from the poles of the black hole, moving away from the viewer.

And the problem is getting even bigger

What to do with all that data?

Why, giant meta-analyses, of course

- ▶ Typical problem across all of big data: how do you use it?
- ▶ In life sciences: no real standards of data formats
- ▶ Data scattered all over, despite push for Data Commons
- ▶ Not always accessible
- ▶ Combining the data if you have it all is a real challenge



A Compounding Problem...

Scientists don't like to share (really!)

- ▶ **The fear:**

- if someone sees data before it is published, they might steal it and publish it themselves (getting scooped)



- ▶ **Causes:**

- Long time to publication
- Outdated methods of assigning scientific credit
- Not properly incentivized



A Problem for Data Commons


Sharing required

- ▶ **Data piling up**
- ▶ **Bad network**



**Hyperscale analytics will only work
if the data is accessible!**

- ▶ **Wild-west file formats/algorithms**
- ▶ **No sharing**





But wait! There's more!

File & Data Types

We have them all

- ▶ **Massive text files**
- ▶ **Massive binary files**
- ▶ **Flatfile ‘databases’**
- ▶ **Spreadsheets everywhere**
- ▶ **Directories w/ 6 million files**
- ▶ **Large files: 600GB+**
- ▶ **Small files: 30kb or smaller**



Application characteristics


- ▶ **Mostly SMP/threaded apps performance bound by IO and/or RAM**
- ▶ **Hundreds of apps, codes & toolkits**
- ▶ **1TB - 2TB RAM “High Memory” applications (large graphs, genomic assembly)**
- ▶ **Lots of Perl/Python/R**
- ▶ **MPI is rare**
 - Well written MPI is even rarer
- ▶ **Few MPI apps actually benefit from expensive low-latency interconnects***
 - *Chemistry, modeling and structure work is the exception



Clinical/Regulated more complicated

Need to protect PII, sensitive/confidential information

- ▶ Historically difficult to do
- ▶ Opposite of the Science DMZ model: compliance at the cost of performance
- ▶ Deeper security than enterprise, mostly because the SSO and compliance officers don't know better (safest?)
- ▶ Doesn't work anymore, data generated from clinics and the field increasing rapidly





Technologies + regulated data + HPC

If IT gets it wrong ...

- ▶ **Lost opportunity**
- ▶ **Missing capability**
- ▶ **Frustrated & very vocal scientific staff**
- ▶ **Problems in recruiting, retention,
publication & product development**

Successful Bio-IT orgs

IT staff become a part of the research projects

- ▶ **When IT and scientists at the table together, mutual understanding happens**
- ▶ **IT is driven to innovate to accommodate science**
- ▶ **Scientists better understand resourcing**
- ▶ **Roadblocks are removed, science is enabled**

Build for the Use Case!

Don't just build for performance or be sold by vendors

- ▶ **Building for hardware is short-sighted**
- ▶ **System needs to solve problems, not create more for scientists**
- ▶ **Fun to engineer for speed, to buy shiny things, but not terribly useful in LifeSci if it doesn't solve a problem**
- ▶ **That means: take the time to understand the use cases, and design to that. Way more fun in the long-term**



Life Scientists: Thugs of HPC


If you have it, they will break it

- ▶ Every kind of file you can imagine
- ▶ Every kind of algorithm (good, bad, ugly)
- ▶ Want it all, right now
- ▶ They keep everything
- ▶ Don't want to pay for it
- ▶ Your health might depend on it (seriously!)



Networking in Life Science

- ▶ **May surpass storage as our #1 infrastructure headache**
- ▶ **Why?**
 - Petascale storage meaningless if you can't access/move it
 - Enterprise network security usually inhibits research data flow
 - 10-Gig, 40-Gig and 100-Gig networking will force significant changes elsewhere in the 'bio-IT' infrastructure







Getting Data out of the Laboratory

Laboratories not Integrated

Usually very little IT infrastructure in labs


- ▶ Tons of data generating equipment going in now
- ▶ Can generate 15GB of data in 50 hours
- ▶ Others can generate 64GB/day
- ▶ Labs are not designed to transmit data, lucky if wired for ethernet



Getting data out

OK, so write data over ethernet to network drive...

- ▶ **Sounds good, 64GB in 24 hours \approx 6Mb/s**
- ▶ **Problem: desktop class ethernet adaptors**
- ▶ **No error checking, no retries, no MD5, no local buffer**
- ▶ **If network goes, whole run is lost**
- ▶ **Also, definitely no encryption**



Getting data out

Clinicians have to get creative, but not in a good way

- ▶ **Usually ends up going to local workstation, locked down by IT**
- ▶ **Go buy the cheapest disks they can, sometimes encrypted USB disks**
- ▶ **Carry it somewhere**
 - Transfer the data to a workstation,
 - Put into compliant database, EHR
 - Not good for analytics
 - HPC use forbidden unless depersonalized, or dedicated HPC



Lab data transit not huge!

Unless you're dealing with a bigger lab with lots of equipment, or a core facility


- ▶ **Fast networking not required, 100Mb OK**
- ▶ **Just GOOD networking**
- ▶ **But!**
 - Collections of data eventually need to be moved for analysis
 - Now you need fast networking
 - Given the problems: how?



Successful models

Some generalized network models that have successfully solved the problem

- ▶ **Most of it is protocol and topology**
- ▶ **Quality of Service (QoS)**
- ▶ **Appropriate segmentation (L2 and/or L3)**
- ▶ **MPLS paths**
- ▶ **Intermediate protocols (i.e., Aspera FASP)**
- ▶ **One way or another, guarantee transfer**
- ▶ **Encrypt the data in flight from data transfer node**





Storing the Data

Storage: a networking problem

As storage needs increase, the need to transmit it goes up too

- ▶ **Networking will quickly replace storage as #1 headache in Bio-IT**
- ▶ **Petascale storage is useless without high-performance networking**
- ▶ **Most enterprise networks won't cut it**

Storing Clinical Data

Huge problem!

- ▶ **Need isolated storage (usually)**
- ▶ **Need data encrypted at rest/in-flight**
- ▶ **Keys need to be stored somewhere else secure**
- ▶ **Keep PII/PHI safe from espionage at all costs**
- ▶ **Usually means you can't use institutional storage**

Storage networking: solutions

Internal networking more critical than external
for petascale storage



- ▶ **Infrastructure must be able to support the inevitable 1PB transit**
 - Disaster recovery
 - High-availability
 - Backup
- ▶ **Need at least 10Gb**
 - Probably dedicated 10Gb per >1PB storage facility: 40Gb min —> 1Tb backbone
- ▶ **1Gb will not cut it for that data size**
 - ~97 days to transmit at saturation
 - 10Gb: ~9.7 days



Storage networking: solutions

And now, the real problem: topology and logical design

- ▶ **Need a scaling internal topology**
- ▶ **One core switch doing all routing and packet transit == bad**
- ▶ **More advanced designs needed**
- ▶ **Also: prioritize performance over security**
 - Nearly impossible for most orgs
- ▶ **Most implemented option: Science DMZ**
- ▶ **Additional implementations for segmentation, isolation, filtering, and encryption: small ISP topology very useful for protecting data**






Analyzing the data

Compute == Answers!


The pinnacle of data transit, the reason we store it in the first place

- ▶ **High performance computing: clusters, supercomputers, single servers, powerful workstations, etc.**
- ▶ **Mostly a datacenter issue**
- ▶ **Unless...**
 - Storage not centralized or co-located: data duplicated unless have a killer network
 - New methods: data doesn't move, compute moves to data





Sharing Scientific data: Collaboration



Collaboration

Fundamental to science

- ▶ Now that data production is reaching petascale, collaboration is getting harder
- ▶ Clinical collaboration is already difficult enough, increased data production in HIPAA environments, larger cohorts, multi-site studies
- ▶ Projects are getting more complex, more data is being generated, takes more people to work on the science
- ▶ Journal authorships: common to see 40+ authors now
- ▶ Clearly a networking problem at its core
- ▶ Let's face it, doing this right is **expensive!**

Data Movement & Data Sharing


The gist of collaborative data sharing in life sciences

- ▶ **Peta-scale data movement needs**

- Within an organization
- To/from collaborators
- To/from suppliers
- To/from public data repos


- ▶ **Peta-scale data sharing needs**

- Collaborators and partners may be all over the world



Network: ‘ScienceDMZ’

- ▶ **Very fast “low-friction” network links and paths with security policy and enforcement specific to scientific workflows**
- ▶ **“ScienceDMZ” concept is real and necessary**
- ▶ **Central premise:**
 - Legacy firewall, network and security methods architected for “many small data flows” use cases
 - **Not built to handle smaller #s of massive data flows**
 - Also very hard to deploy ‘traditional’ security gear on 10Gigabit and faster networks
- ▶ **More details, background & documents at**
<http://fasterdata.es.net/science-dmz/>



Simple Science DMZ:






Image source: “The Science DMZ: Introduction & Architecture” -- esnet


Standard Science DMZ




Trusted/Untrusted Science DMZ



Medical/Regulated Science DMZ



Multi-site Science DMZ





What's the Big Picture?


[Hyper-]convergence

Make IT Infrastructure a research tool



Laboratory



Converged Solution




Knowledge



Universal Truth

“The network IS the computer” - *John Gage, Sun Microsystems*

- ▶ **Convergence is not possible without networking**
- ▶ **Also not possible without GOOD networking**
- ▶ **Life Sciences is learning lessons learned by physics and astronomy 5-10 years ago**
- ▶ **Biggest problem is Org acceptance and investment in personnel and equipment**
- ▶ **Next-Gen biomedical research advancing too quickly: must invest now**





end; Thanks!